

Contact

www.linkedin.com/in/bhomass
(LinkedIn)

Top Skills

Software Development
Java
SOA

Certifications

AWS Certified Solutions Architect - Associate

Publications

A real time streaming
implementation of markov chain
based fraud detection
Introduction to Time Series
Multiple RDDs write to S3 in parallel

Bruce Ho

Chief Data Scientist
Orange County, California Area

Summary

Top rated problem solver and communicator. Unique strength in bridging software engineering with mathematical algorithms. Advanced education and research training at MIT and Caltech. Specialty in Cloud solutions, Big Data engineering and analytics. Over 15 years IT experience in both high-visibility large-scale projects as well as special targeted problem solving. Breadth in both application domain and technology stacks.

- Big data analytics and engineering, real time streaming apps using Spark
- Advanced Spark techniques, partition strategy, maximizing parallelization, RDD reuse, YARN optimization
- Machine learning / predictive modeling / TB data pipeline / time critical data ingestion
- Experience with a host of NoSQL databases, Cassandra schema, performance benchmarking
- Expertise in AWS cloud environment, EMR, Kinesis, S3, SQS, RDS, Redshift, VPC, etc.
- Well versed in math techniques of machine learning. Proposed original algorithm for digital ad attribution modeling. Direct translation of high-level business goals into machine learning project implementation
- Architectural and leadership role for the past 10+ years, mostly as direct report to CTO and VPs.
- Multi-tier enterprise scale java application development and multi-center integration using service oriented architecture (SOA)
- Managed development team both overseas and in the States. Familiar with outsourcing in India, Bangkok and Manila
- Published over 100 articles in research journals, book chapters, and conference proceedings, principal investigator of a federal grant

Experience

UST Global

Resident Data Scientist

April 2018 - Present (2 years)

Aliso Viejo, CA

nlp, entity relation extraction, topic modeling, automated question & answer
speech recognition, action item extraction, audio analysis
image recognition, retail application, on shelf availability, deep instance
retrieval, face recognition

BigR.io

Chief Data Scientist

April 2016 - Present (4 years)

Creating enterprise data hubs for client companies. All things Big Data, data lake, streaming analytics, EDW, NoSQL, AWS cloud solutions. Challenging Machine Learning problems, HMM, Hierarchical Bayesian, MCMC, Deep Learning, behavior modeling, IoT, TensorFlow, and t-SNE.

BigR.io is a technology consulting firm empowering data to drive innovation and advanced analytics. We specialize in cutting-edge Big Data, Machine Learning, and Custom Software strategy, analysis, architecture, and implementation solutions. We are an elite group with MIT roots, shining when tasked with complex missions. Whether it's assembling mounds of data from a variety of sources, surfacing intelligence with Deep Learning, or building high-volume, highly-available systems, we consistently deliver.

With extensive domain knowledge, BigR.io's scientists and engineers design and build best-in-class solutions across a variety of verticals. This diverse industry exposure and our constant run-in with cutting-edge technology equips us with invaluable tools, tricks, and techniques. Our knowledge and horsepower bring innovative, cost-conscious, and extensible results to complex business challenges.

Clouddiy, Inc.

CTO

2005 - Present (15 years)

Offer mobile app IDE with super streamlined user interface that business analysts can easily create and deploy business apps.

OrangePeople

Advisory Consultant

June 2016 - Present (3 years 10 months)

IoT strategy, advanced predictive analytics in support of business decision making.

CAKE (getCAKE.com)

Big Data Machine Learning Architect

May 2015 - April 2016 (1 year)

Newport Beach, CA

Big Data platform, Spark, Cassandra, Kafka, Redshift, DynamoDB, Kinesis, predicative modeling, ad attribution, resource allocation optimization.

Adaptive Medias, Inc

Senior Software Architect

May 2014 - May 2015 (1 year 1 month)

Irvine CA

Big Data Analytics, Machine learning, Spark, Cassandra, Hbase, Hadoop, Scala.

Amazon Lab126

Senior Software Developer

January 2013 - July 2013 (7 months)

Sunnyvale, CA

Key developer on Kindle ticket servicing system and mobile survey app.

Dell Services

Enterprise Architect

September 2011 - September 2012 (1 year 1 month)

Lake Forest, CA

Frameworks, Integration, System maintenance, Technology research.

Teradata

Java Architect

January 2011 - August 2011 (8 months)

Created a GWT based expression canvas for generating data mining queries.

Harvard Medical School

Faculty, Medical Imaging

1996 - 1999 (3 years)

Education

Massachusetts Institute of Technology

BS, Physics

Caltech

PhD, Applied Physics